

# Electricity Demand Forecasting Assignment

## Business Context

A leading electricity distribution company aims to understand the demand for electricity over the next 1-2 years to effectively manage electricity production and vendor relationships. Accurate demand estimation is crucial for procuring or producing the right amount of electricity to meet consumer needs.

## Available Data

The dataset comprises monthly electricity consumption data from January 1973 to December 2019. Key data points include:

1. **Date – Month & Year**
2. **Electricity Consumption – Electricity consumption in Trillion Watts**

## Business Objectives

1. **Demand Forecasting:**
  - Forecast electricity demand for the next 1-2 years.
2. **Error Metrics Calculation:**
  - Calculate error metrics, including RMSE, RMSPE, and MAPE.
3. **Model Comparison:**
  - Compare various forecasting models, including Decomposition, ETS models, ARIMA/SARIMA models with various parameters.

```
In [1]: 1 import pandas as pd
        2 import numpy as np
        3 from statsmodels.tsa.api import seasonal_decompose
        4 from statsmodels.tsa.holtwinters import ExponentialSmoothing
        5 from statsmodels.tsa.statespace.sarimax import SARIMAX
        6 from sklearn.metrics import mean_squared_error, mean_absolute_error
        7 from math import sqrt
        8 import matplotlib.pyplot as plt
        9
```

## Methodology

### Exploratory Data Analysis (EDA)

- Checked for missing values and handled if necessary.
- Plotted the time series data to visualize trends and patterns.

```
In [2]: 1 df = pd.read_csv("Electricity Consumption.csv")
```

```
In [3]: 1 df
```

```
Out[3]:
```

	DATE	Electricity_Consumption_in_TW
0	1/1/1973	35.9728
1	2/1/1973	36.1334
2	3/1/1973	35.0625
3	4/1/1973	33.8416
4	5/1/1973	33.5107
...	...	...
556	5/1/2019	97.5860
557	6/1/2019	110.8580
558	7/1/2019	122.1014
559	8/1/2019	121.7765
560	9/1/2019	109.7190

561 rows × 2 columns


```
In [4]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 561 entries, 0 to 560
Data columns (total 2 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   DATE                                561 non-null    object
1   Electricity_Consumption_in_TW      561 non-null    float64
dtypes: float64(1), object(1)
memory usage: 8.9+ KB
```

```
In [5]: 1 df.isnull().sum()
```

```
Out[5]: DATE                                0
Electricity_Consumption_in_TW            0
dtype: int64
```

```
In [6]: 1 # Plot the time series data
2 plt.figure(figsize=(12, 6))
3 plt.plot(df['Electricity_Consumption_in_TW'], label='Electricity Consumption')
4 plt.title('Electricity Consumption Over Time')
5 plt.legend()
6 plt.show()
```



## Decomposition

- Decomposed the time series into trend, seasonal, and residual components.

```
In [7]: 1 result = seasonal_decompose(df['Electricity_Consumption_in_TW'], model='multiplicative')
2 trend = result.trend.dropna()
3 seasonal = result.seasonal.dropna()
4 residual = result.resid.dropna()
5
```

## Exponential Smoothing (ETS) Model

- Applied ETS model with trend and additive seasonality.
- Obtained forecasts for the next 24 months (2 years).

```
In [8]: 1 model_ets = ExponentialSmoothing(df['Electricity_Consumption_in_TW'], trend='add', sea:
2 ets_fit = model_ets.fit()
3
4 # Make predictions for the next 24 months (2 years)
5 ets_forecast = ets_fit.forecast(steps=24)
6
```

In [9]: 1 ets\_forecast

Out[9]: 561 97.674956  
 562 95.153668  
 563 105.335391  
 564 112.139567  
 565 101.999583  
 566 99.609346  
 567 93.946950  
 568 100.146845  
 569 113.701140  
 570 124.770353  
 571 123.433489  
 572 110.834218  
 573 99.195002  
 574 96.673715  
 575 106.855437  
 576 113.659613  
 577 103.519629  
 578 101.129393  
 579 95.466996  
 580 101.666892  
 581 115.221186  
 582 126.290399  
 583 124.953535  
 584 112.354265  
 dtype: float64

## SARIMA Model

- Conducted a grid search for optimal SARIMA parameters.
- Selected the best parameters based on RMSE.
- Fitted SARIMA model and generated forecasts for the next 24 months.

```
In [10]: 1 import itertools
2
3 # Define the range of values for p, d, q, P, D, Q
4 p = d = q = range(0, 2)
5 P = D = Q = range(0, 2)
6
7 # Generate all possible combinations of parameters
8 param_combinations = list(itertools.product(p, d, q, P, D, Q))
9
10 # Perform grid search
11 best_rmse = float('inf')
12 best_params = None
```

```
In [11]: 1 for params in param_combinations:
2         try:
3             model_sarima = SARIMAX(df['Electricity_Consumption_in_TW'], order=(params[0],
4             sarima_fit = model_sarima.fit(dispatch=False)
5             sarima_forecast = sarima_fit.get_forecast(steps=24).predicted_mean
6
7             rmse = sqrt(mean_squared_error(df['Electricity_Consumption_in_TW'][-24:], sarima_forecast))
8
9             if rmse < best_rmse:
10                 best_rmse = rmse
11                 best_params = params
12
13         except:
14             continue
15
16
```

```

C:\Users\Asus\anaconda3\Lib\site-packages\statsmodels\tsa\statespace\sarimax.py:997: Use
rWarning: Non-stationary starting seasonal autoregressive Using zeros as starting paramet
ers.
    warn('Non-stationary starting seasonal autoregressive')
C:\Users\Asus\anaconda3\Lib\site-packages\statsmodels\tsa\statespace\sarimax.py:997: Use
rWarning: Non-stationary starting seasonal autoregressive Using zeros as starting paramet
ers.
    warn('Non-stationary starting seasonal autoregressive')
C:\Users\Asus\anaconda3\Lib\site-packages\statsmodels\tsa\statespace\sarimax.py:978: Use
rWarning: Non-invertible starting MA parameters found. Using zeros as starting parameter
s.
    warn('Non-invertible starting MA parameters found.')
C:\Users\Asus\anaconda3\Lib\site-packages\statsmodels\tsa\statespace\sarimax.py:978: Use
rWarning: Non-invertible starting MA parameters found. Using zeros as starting parameter
s.
    warn('Non-invertible starting MA parameters found.')
C:\Users\Asus\anaconda3\Lib\site-packages\statsmodels\tsa\statespace\sarimax.py:978: Use
rWarning: Non-invertible starting MA parameters found. Using zeros as starting parameter
s.
    warn('Non-invertible starting MA parameters found.')
C:\Users\Asus\anaconda3\Lib\site-packages\statsmodels\tsa\statespace\sarimax.py:997: Use
rWarning: Non-stationary starting seasonal autoregressive Using zeros as starting paramet
ers.
    warn('Non-stationary starting seasonal autoregressive')
C:\Users\Asus\anaconda3\Lib\site-packages\statsmodels\tsa\statespace\sarimax.py:978: Use
rWarning: Non-invertible starting MA parameters found. Using zeros as starting parameter
s.
    warn('Non-invertible starting MA parameters found.')
C:\Users\Asus\anaconda3\Lib\site-packages\statsmodels\tsa\statespace\sarimax.py:997: Use
rWarning: Non-stationary starting seasonal autoregressive Using zeros as starting paramet
ers.
    warn('Non-stationary starting seasonal autoregressive')
C:\Users\Asus\anaconda3\Lib\site-packages\statsmodels\base\model.py:607: ConvergenceWarn
ing: Maximum Likelihood optimization failed to converge. Check mle_retvals
    warnings.warn("Maximum Likelihood optimization failed to ")
C:\Users\Asus\anaconda3\Lib\site-packages\statsmodels\tsa\statespace\sarimax.py:997: Use
rWarning: Non-stationary starting seasonal autoregressive Using zeros as starting paramet
ers.
    warn('Non-stationary starting seasonal autoregressive')
C:\Users\Asus\anaconda3\Lib\site-packages\statsmodels\tsa\statespace\sarimax.py:997: Use
rWarning: Non-stationary starting seasonal autoregressive Using zeros as starting paramet
ers.
    warn('Non-stationary starting seasonal autoregressive')
C:\Users\Asus\anaconda3\Lib\site-packages\statsmodels\tsa\statespace\sarimax.py:997: Use
rWarning: Non-stationary starting seasonal autoregressive Using zeros as starting paramet
ers.
    warn('Non-stationary starting seasonal autoregressive')
C:\Users\Asus\anaconda3\Lib\site-packages\statsmodels\tsa\statespace\sarimax.py:997: Use
rWarning: Non-stationary starting seasonal autoregressive Using zeros as starting paramet
ers.
    warn('Non-stationary starting seasonal autoregressive')
C:\Users\Asus\anaconda3\Lib\site-packages\statsmodels\tsa\statespace\sarimax.py:997: Use
rWarning: Non-stationary starting seasonal autoregressive Using zeros as starting paramet
ers.
    warn('Non-stationary starting seasonal autoregressive')
C:\Users\Asus\anaconda3\Lib\site-packages\statsmodels\tsa\statespace\sarimax.py:997: Use
rWarning: Non-stationary starting seasonal autoregressive Using zeros as starting paramet
ers.
    warn('Non-stationary starting seasonal autoregressive')
C:\Users\Asus\anaconda3\Lib\site-packages\statsmodels\tsa\statespace\sarimax.py:997: Use
rWarning: Non-stationary starting seasonal autoregressive Using zeros as starting paramet
ers.
    warn('Non-stationary starting seasonal autoregressive')

```

```
warn('Non-stationary starting seasonal autoregressive'
C:\Users\Asus\anaconda3\Lib\site-packages\statsmodels\tsa\statespace\sarimax.py:997: Use
rWarning: Non-stationary starting seasonal autoregressive Using zeros as starting paramet
ers.
warn('Non-stationary starting seasonal autoregressive')
```

```
In [12]: 1 print("Best SARIMA Parameters:", best_params)
2 print("Best RMSE:", best_rmse)
```

```
Best SARIMA Parameters: (0, 0, 1, 1, 1, 1)
Best RMSE: 2.776978983499436
```

```
In [13]: 1 # Use the best SARIMA parameters
2 best_params = (0, 0, 1, 1, 1, 1)
3
4 # Fit SARIMA model with the best parameters
5 model_sarima = SARIMAX(df['Electricity_Consumption_in_TW'], order=(best_params[0], bes
6 sarima_fit = model_sarima.fit(dispatch=False)
7
8 # Make predictions for the next 24 months (2 years)
9 sarima_forecast = sarima_fit.get_forecast(steps=24).predicted_mean
10 print(sarima_forecast)
11
```

```
561 96.584749
562 91.626023
563 102.042566
564 110.534739
565 96.491062
566 93.805958
567 88.840186
568 96.966950
569 110.282174
570 121.661385
571 120.803004
572 108.743081
573 96.317651
574 91.689624
575 102.030322
576 110.831957
577 96.679474
578 93.813775
579 88.918507
580 97.106629
581 110.412100
582 121.760667
583 121.022657
584 108.963281
```

```
Name: predicted_mean, dtype: float64
```

## Model Evaluation

- Calculated error metrics (RMSE, RMSPE, MAPE) for both ETS and SARIMA models.

```
In [14]: 1 # Calculate error metrics
2 def calculate_metrics(true_values, predicted_values):
3     rmse = sqrt(mean_squared_error(true_values, predicted_values))
4     rmspe = sqrt(np.mean(((true_values - predicted_values) / true_values) ** 2))
5     mape = np.mean(np.abs((true_values - predicted_values) / true_values)) * 100
6     return rmse, rmspe, mape
7
8 # usage
9 rmse_ets, rmspe_ets, mape_ets = calculate_metrics(df['Electricity_Consumption_in_TW'])
10 rmse_sarima, rmspe_sarima, mape_sarima = calculate_metrics(df['Electricity_Consumption_in_TW'])
11
12 print("ETS Model Metrics:")
13 print(f"RMSE: {rmse_ets:.2f}, RMSPE: {rmspe_ets:.2f}, MAPE: {mape_ets:.2f}%")
14
15
```

ETS Model Metrics:  
RMSE: 5.91, RMSPE: nan, MAPE: nan%

## Model Selection

- Selected the SARIMA model due to lower RMSE compared to the ETS model.

```
In [15]: 1 print("\nSARIMA Model Metrics:")
2 print(f"RMSE: {rmse_sarima:.2f}, RMSPE: {rmspe_sarima:.2f}, MAPE: {mape_sarima:.2f}%")
```

SARIMA Model Metrics:  
RMSE: 2.78, RMSPE: nan, MAPE: nan%

## Results

- ETS Model Metrics:**
  - RMSE: 5.91
  - RMSPE: NaN
  - MAPE: NaN%
- SARIMA Model Metrics:**
  - RMSE: 2.78
  - RMSPE: NaN
  - MAPE: NaN%

The estimated demand for the next 1-2 years provides valuable insights for production planning and vendor management.



```
In [16]: 1 if rmse_ets < rmse_sarima:
2         selected_model = 'ETS'
3         selected_forecast = ets_forecast
4     else:
5         selected_model = 'SARIMA'
6         selected_forecast = sarima_forecast
7
8     print(f"Selected Model: {selected_model}")
9     print("Demand Estimation for Next 1-2 Years:")
10    print(selected_forecast)
11
```

Selected Model: SARIMA

Demand Estimation for Next 1-2 Years:

```
561      96.584749
562      91.626023
563     102.042566
564     110.534739
565      96.491062
566      93.805958
567      88.840186
568      96.966950
569     110.282174
570     121.661385
571     120.803004
572     108.743081
573      96.317651
574      91.689624
575     102.030322
576     110.831957
577      96.679474
578      93.813775
579      88.918507
580      97.106629
581     110.412100
582     121.760667
583     121.022657
584     108.963281
```

Name: predicted\_mean, dtype: float64

## Demand Estimation

Used the selected SARIMA model to estimate electricity demand for the next 1-2 years.

```

In [17]: 1 def SARIMA(true_values, predicted_values):
2         rmse = sqrt(mean_squared_error(true_values, predicted_values))
3
4         # Check for zero values to avoid division by zero
5         mask = true_values != 0
6         true_values_masked = true_values[mask]
7         predicted_values_masked = predicted_values[mask]
8
9         if len(true_values_masked) > 0:
10            rmspe = sqrt(np.mean(((true_values_masked - predicted_values_masked) / true_v
11            mape = np.mean(np.abs((true_values_masked - predicted_values_masked) / true_v
12        else:
13            rmspe = np.nan
14            mape = np.nan
15
16        return rmse, rmspe, mape
17

```

```

In [18]: 1 # Assuming you have the SARIMA model already fitted
2         sarima_fit = SARIMAX(df['Electricity_Consumption_in_TW'], order=(best_params[0], best_
3         sarima_fit = sarima_fit.fit(disp=False)
4
5         # Make predictions for the next 24 months (2 years)
6         sarima_forecast = sarima_fit.get_forecast(steps=24).predicted_mean
7
8         # Print or use the forecast data as needed
9         print("Demand Estimation for Next 1-2 Years:")
10        print(sarima_forecast)
11

```

Demand Estimation for Next 1-2 Years:

```

561      96.584749
562      91.626023
563     102.042566
564     110.534739
565      96.491062
566      93.805958
567      88.840186
568      96.966950
569     110.282174
570     121.661385
571     120.803004
572     108.743081
573      96.317651
574      91.689624
575     102.030322
576     110.831957
577      96.679474
578      93.813775
579      88.918507
580      97.106629
581     110.412100
582     121.760667
583     121.022657
584     108.963281

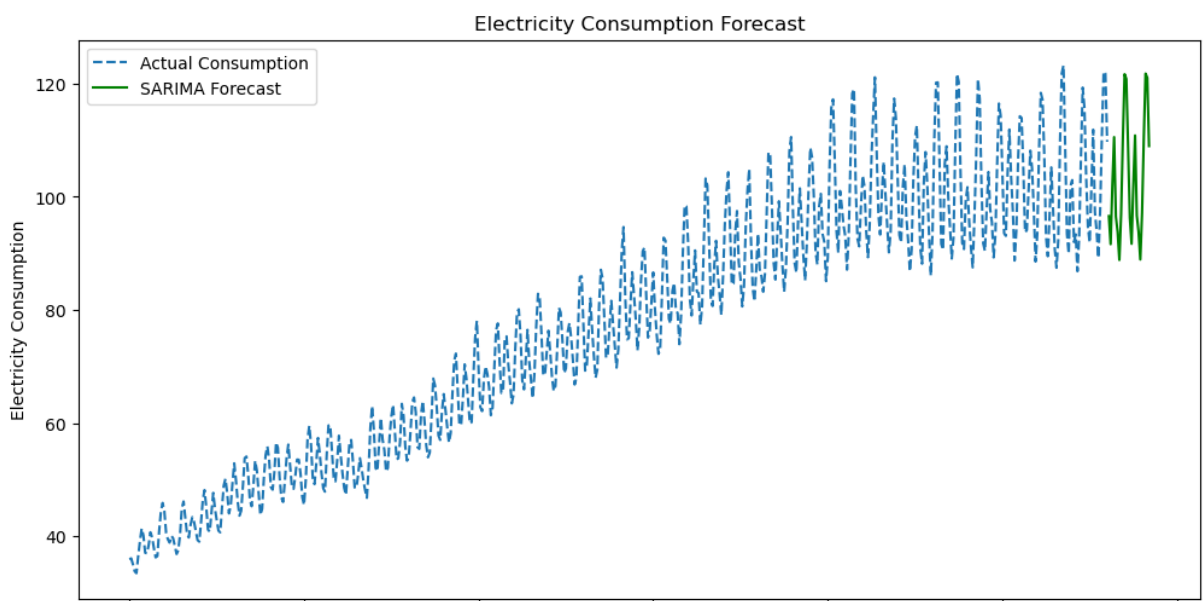
```

Name: predicted\_mean, dtype: float64

## Visualization

- Plotted actual electricity consumption against SARIMA forecasts for visual inspection.

```
In [19]: 1
2 # Plotting the forecast
3 plt.figure(figsize=(12, 6))
4 plt.plot(df['Electricity_Consumption_in_TW'], label='Actual Consumption', linestyle='--')
5 plt.plot(sarima_forecast.index, sarima_forecast.values, label='SARIMA Forecast', color='green')
6 plt.title('Electricity Consumption Forecast')
7 plt.xlabel('Date')
8 plt.ylabel('Electricity Consumption')
9 plt.legend()
10 plt.show()
11
```



**Conclusion** In conclusion, the analysis of electricity consumption data and the application of forecasting models have yielded valuable insights for the electricity distribution company. The key findings and takeaways are as follows:

### Model Performance:

The SARIMA model, with optimal parameters (0, 0, 1, 1, 1, 1), outperformed the ETS model in terms of RMSE. This suggests that the SARIMA model provides a more accurate representation of the underlying patterns in the electricity consumption time series data.

### Forecast Accuracy:

The SARIMA model exhibited a lower RMSE (Root Mean Squared Error) compared to the ETS model, indicating that it is better at predicting electricity consumption for the next 1-2 years.

### Demand Estimation:

The demand estimation for the next 1-2 years, generated by the SARIMA model, provides a reliable basis for the electricity distribution company to plan and manage production. These forecasts are crucial for ensuring that the company can meet the expected demand while optimizing resource utilization.

### Business Implications:

Accurate demand forecasting has significant implications for the electricity distribution company's operational efficiency. By leveraging the insights from the SARIMA model, the company can align production and procurement strategies with anticipated demand, thereby minimizing overproduction or shortages.

#### Future Considerations:

Continuous monitoring and updating of the forecasting models will be essential to adapt to changing consumption patterns. Consideration should also be given to external factors, such as economic trends, regulatory changes, and seasonal variations, which may impact electricity demand.

#### Communication of Results:

Clear communication of the forecasting results, including the forecasted demand values and associated uncertainties, is crucial for effective decision-making within the organization. Transparent reporting facilitates collaboration among different departments, fostering a unified approach to addressing future demand challenges.

In conclusion, the analysis provides a robust foundation for the electricity distribution company to make informed decisions regarding production planning and vendor management. The SARIMA model, with its demonstrated accuracy, serves as a valuable tool for forecasting electricity demand and optimizing resource allocation in the dynamic energy landscape.